


Generative AI + Law

ICML: July 29, 2023



Goals

Precision

Common language

Shared research directions

Time	Speaker
9:15 - 9:45 AM	IP
9:45 - 10:05 AM	
10:05 - 10:40 AM	
10:40 - 11:00 AM	
11:00 - 12:00 PM	
12:00 - 1:30 PM	Spotlights + Posters
1:30 - 1:45 PM	
1:45 - 2:15 PM	Privacy
2:15 - 3:00 PM	
3:00 - 3:30 PM	
3:30 - 3:45 PM	
3:45 - 4:00 PM	
4:00 - 5:00 PM	

Generative AI + Law



Time	Speaker
9:15 - 9:45 AM	Pamela Samuelson , “Some Nonobvious Observation About Copyright’s Scope For Generative AI Developers”
9:45 - 10:05 AM	Mark Lemley , “Is Training AI Copyright Infringement?”
10:05 - 10:40 AM	Coffee
10:40 - 11:00 AM	Miles Brundage , “Where and when does the law fit into AI development and deployment?”
11:00 - 12:00 PM	<i>Panel on Intellectual Property:</i> Pamela Samuelson , Mark Lemley , Luis Villa , Katherine Lee (Moderated by Jack Balkin and A. Feder Cooper)
12:00 - 1:30 PM	Lunch
1:30 - 1:45 PM	Jack Balkin , “Artificial Intelligence and the First Amendment”
1:45 - 2:15 PM	Spotlights
2:15 - 3:00 PM	Posters
3:00 - 3:30 PM	Coffee
3:30 - 3:45 PM	Nicholas Carlini , “A Brief Introduction to Machine Learning & Memorization”
3:45 - 4:00 PM	Gautam Kamath , “What does Differential Privacy have to do with Copyright?”
4:00 - 5:00 PM	<i>Panel on Privacy:</i> Kristen Vaccaro , Nicholas Carlini , Miles Brundage , Gautam Kamath , and Jack Balkin (Moderated by Katherine Lee and Deep Ganguli)

Generative AI + Law



Speakers, Panelists & Moderators



Pamela Samuelson

*Distinguished Professor of
Law and Information*
University of California,
Berkeley



Mark Lemley

Professor of Law
Stanford Law School



Nicholas Carlini

Research Scientist
Google Brain



Gautam Kamath

Assistant Professor
University of Waterloo



Kristen Vaccaro

Assistant Professor
University of California, San
Diego



Luis Villa

*Co-founder and General
Counsel*
Tidelift



Miles Brundage

Head of Policy Research
OpenAI



Jack M. Balkin

Professor
Yale Law School

What are we talkin' about?

Dataset creation

Model training

Model serving

Dataset creation

Model training

Model serving

Training **dataset**
creation involves a
lot of *choices*

What type of data?
How much data?

What type of data?

How much data?

Web data?

Which web data?

What type of data?

How much data?

Web data?

Which web data?

“Good” data?

What type of data?

How much data?

Web data?

Which web data?

“Good” data?

Copyrighted data?

Private data?

Training dataset creation involves a lot of *choices*

The Devil is In the Training Data

Chapter 1

AUTHORS

Katherine Lee


Daphne Ippolito

A. Feder Cooper

PUBLISHED

July 10, 2023

The process of training contemporary generative models requires vast quantities of *training data*. Dataset creators and curators make extensive decisions about how much and which data to include in a training dataset. These choices directly and significantly shape a model's outputs (a.k.a. *generations*), including the model's capacity to learn concepts and produce novel content.




<https://genlaw.github.io/explainers/>

Dataset creation

Model training

Model serving

Model Training



Model **training**
involves a lot of
choices

Dataset creation
Model training
Model serving

Dataset creation

Model training

Model serving

Prompting

Generation

Fine-tuning

Human Feedback

Prompting & Generation

Prompt

Two little robots sitting
on the beach reading
law books




Model




Generation




Fine-tuning



Fine-tuning



Human Feedback



Model **serving**
involves a lot of
choices

Choices

Choices



Research Questions

(1) GETTY IMAGES (US), INC.

(a company incorporated under the laws of the State of New York)

(2) GETTY IMAGES INTERNATIONAL J.L., C.B., K.S., P.M., N.G., R.F., J.D. and G.D.

(a company incorporated under the laws of individually, and on behalf of all others PAUL TREMBLAY, an individual and

(3) GETTY IMAGES (UK) LIMITED situated, MONA AWAD, an individual,

(4) GETTY IMAGES DEVCO UK LIMIT

Individual and Representative Plaintiffs

SARAH ANDERSEN, an individual;
KELLY MCKERNAN, an individual;
KARLA ORTIZ, an individual,

Plaintiffs:
RAH SILVERMAN, an individual;
RISTOPHER GOLDEN, an individual;
CHARD KADREY, an individual;

Individual and Representative Plaintiffs,

Individual and Representative Plaintiffs,

v.

v.

STABILITY AI LTD., a UK corporation;
STABILITY AI, INC., a Delaware
corporation; MIDJOURNEY, INC., a
Delaware corporation; DEVIANTART, INC.,
a Delaware corporation,

Defendants.

v.

META PLATFORMS, INC., a Delaware

ENAI, INC., a Delaware nonprofit corporation; OPENAI, L.P., a
aware limited partnership; OPENAI OPco, L.L.C., a Delaware
ted liability corporation; OPENAI GP, L.L.C., a Delaware
ted liability company; OPENAI STARTUP FUND GP I, L.L.C.,
limited liability company; OPENAI STARTUP FUND I, VS.
aware limited partnership; and OPENAI STARTUP FUND
MENT, LLC, a Delaware limited liability company,

Defendants.

corporation; OPENAI, INC., a Delaware OPENAI STARTUP FUND MAN
profit corporation; OPENAI, L.P., a LLC, MICROSOFT CORPORATI
ware limited partnership; OPENAI 1 through 20, inclusive,
C., a Delaware limited liability com Defendants.
INAI STARTUP FUND GP I, II

Choices



Research Questions

IP + Privacy

Organizers



Katherine Lee

Ph.D. Candidate

Cornell University Department
of Computer Science



A. Feder Cooper

Ph.D. Candidate

Cornell University Department
of Computer Science



**Fatemehsadat
Mireshghallah**

Post-Doctoral Researcher
University of Washington



Madiha Z. Choksi

Ph.D. Student

Cornell University Department
of Information Science



James Grimmelmann

*Professor of Digital and
Information Law*

Cornell Law School and
Cornell Tech



David Mimno

Associate Professor

Cornell University Department
of Information Science



Deep Ganguli

Research Scientist
Anthropic



Ludwig Schubert

Organizers



Katherine Lee

Ph.D. Candidate

Cornell University Department
of Computer Science



A. Feder Cooper

Ph.D. Candidate

Cornell University Department
of Computer Science



**Fatemehsadat
Mireshghallah**

Post-Doctoral Researcher
University of Washington



Madiha Z. Choksi

Ph.D. Student

Cornell University Department
of Information Science



James Grimmelmann

*Professor of Digital and
Information Law*

Cornell Law School and
Cornell Tech



David Mimno

Associate Professor

Cornell University Department
of Information Science



Deep Ganguli

Research Scientist
Anthropic



Ludwig Schubert



Microsoft



ML Collective

Time	Speaker
9:15 - 9:45 AM	Pamela Samuelson , “Some Nonobvious Observation About Copyright’s Scope For Generative AI Developers”
9:45 - 10:05 AM	Mark Lemley , “Is Training AI Copyright Infringement?”
10:05 - 10:40 AM	Coffee
10:40 - 11:00 AM	Miles Brundage , “Where and when does the law fit into AI development and deployment?”
11:00 - 12:00 PM	<i>Panel on Intellectual Property:</i> Pamela Samuelson , Mark Lemley , Luis Villa , Katherine Lee (Moderated by Jack Balkin and A. Feder Cooper)
12:00 - 1:30 PM	Lunch
1:30 - 1:45 PM	Jack Balkin , “Artificial Intelligence and the First Amendment”
1:45 - 2:15 PM	Spotlights
2:15 - 3:00 PM	Posters
3:00 - 3:30 PM	Coffee
3:30 - 3:45 PM	Nicholas Carlini , “A Brief Introduction to Machine Learning & Memorization”
3:45 - 4:00 PM	Gautam Kamath , “What does Differential Privacy have to do with Copyright?”
4:00 - 5:00 PM	<i>Panel on Privacy:</i> Kristen Vaccaro , Nicholas Carlini , Miles Brundage , Gautam Kamath , and Jack Balkin (Moderated by Katherine Lee and Deep Ganguli)

Generative AI + Law



Speakers:



Pamela Samuelson

*Distinguished Professor of
Law and Information*
University of California,
Berkeley



Mark Lemley

Professor of Law
Stanford Law School



Luis Villa

*Co-founder and General
Counsel*
Tidelift



Katherine Lee

Ph.D. Candidate
Cornell University Department
of Computer Science

Moderators:



Jack M. Balkin

Professor
Yale Law School



A. Feder Cooper

Ph.D. Candidate
Cornell University Department
of Computer Science

Speakers:



Kristen Vaccaro

Assistant Professor
University of California, San
Diego



Nicholas Carlini

Research Scientist
Google DeepMind



Gautam Kamath

Assistant Professor
University of Waterloo



Jack M. Balkin

Professor
Yale Law School



Miles Brundage

Head of Policy Research
OpenAI

Moderators:



Katherine Lee

Ph.D. Candidate
Cornell University Department
of Computer Science



Deep Ganguli

Research Scientist
Anthropic

Takeaways

Interpretation

Evaluation

Alignment

Takeaways

Interpretation
Evaluation
Alignment

IP + Privacy